









Intro to Deep Generative Models - DIRP Spring 2025

– Meeting 1: Feb 8, 2025 –
Introduction and Computer Vision Basics

Introductions

Topic	Resources
Computer Vision Basics <ul style="list-style-type: none"> - Image formation and representation - Basic image processing (noise, filtering, denoising) - Convolutions - Edge detection and feature extraction - Introduction to Python libraries (Numpy, OpenCV) 	<ul style="list-style-type: none"> - Image Representation: <ul style="list-style-type: none"> ▶ Image Representation - Linear Filtering: ▶ Linear Filtering - Edge Detection: ▶ Edge Detection - ★ Canny Edge Detector Paper: Computational Approach To Edge Detection - Feature detection: <ul style="list-style-type: none"> ▶ Feature Detectors : SIFT and Variants - 📖 Szeliski Chapter 3.1-3.3
Deep Learning for Vision <ul style="list-style-type: none"> - Basics of a neural network - Backpropagation - Activation functions - Loss functions and optimization 	First three videos of this playlist (3B1B) Neural networks - YouTube
Convolutional Neural Networks (CNNs) <ul style="list-style-type: none"> - CNN Architecture - Convolution and pooling - Popular CNN architectures - Batch normalization - Data augmentation 	<ul style="list-style-type: none"> - What is a convolution? <ul style="list-style-type: none"> ▶ But what is a convolution? - 📖 Deep Learning Book Chapter 9: https://www.deeplearningbook.org/contents/convnets.html
Recurrent Neural Networks (RNNs) <ul style="list-style-type: none"> - RNN architecture - LSTM and GRU 	<ul style="list-style-type: none"> - 📖 Deep Learning Book Chapter 10: https://www.deeplearningbook.org/contents/rnn.html

Convolutional Neural Networks (CNNs) <ul style="list-style-type: none"> - CNN Architecture - Convolution and pooling - Popular CNN architectures - Batch normalization - Data augmentation 	<ul style="list-style-type: none"> - What is a convolution?  But what is a convolution? -  Deep Learning Book Chapter 9: https://www.deeplearningbook.org/contents/convnets.html
Recurrent Neural Networks (RNNs) <ul style="list-style-type: none"> - RNN architecture - LSTM and GRU 	<ul style="list-style-type: none"> -  Deep Learning Book Chapter 10: https://www.deeplearningbook.org/contents/rnn.html
Autoencoders <ul style="list-style-type: none"> - Autoencoder architectures - Latent space/dimensionality reduction - Feature learning - Variational Autoencoders (VAEs) <ul style="list-style-type: none"> - VAEGAN, VQ-VAE 	<ul style="list-style-type: none"> -  Variational Autoencoders → Highly recommend this video - An Introduction to VAE-GANs - ★ Autoencoding beyond pixels paper: https://arxiv.org/abs/1512.09300
Generative Adversarial Networks (GANs) <ul style="list-style-type: none"> - GAN architecture - DCGANs, Conditional GANs 	<ul style="list-style-type: none"> - Friendly Introduction to GANs:  A Friendly Introduction to Generative Ad... - GAN implementation: Build a GAN From Scratch
Advanced Generative Models <ul style="list-style-type: none"> - Score based and energy based models - Diffusion Models 	<ul style="list-style-type: none"> - Energy-Based Models  Stanford CS236: Deep Generative Models I 2023 ...  Stanford CS236: Deep Generative Models I 2023 ... <ul style="list-style-type: none"> - Score-Based Models  Stanford CS236: Deep Generative Models I 2023 ...

Introduction to Deep Generative Models

This is a reading group under the Directed Reading Program (DiRP) at UT Austin. This group will cover concepts starting from basic computer vision concepts such as filtering, convolutions, edge/object detection, etc. and will slowly build up to more complex models including CNNs, RNNs, Variational Autoencoders (VAEs, VAE-GANs, VQ-VAEs), GANs, and Diffusion Models. This website is intended to be a central hub for notes, papers, and many other resources that will be helpful to anyone researching this topic.


Website for course notes + other resources!

What is an image? How do we represent it?



- Images are 2D representations of a 3D scene
- Three major ways of representing images:
 - Matrix form
 - Function representation
 - Digital images

Matrix Representation



Digitized "A"

0	0	1	1	0	0
0	0	1	1	0	0
0	1	1	1	1	0
0	1	0	0	1	0
0	1	1	1	1	0
1	1	0	0	1	1
1	0	0	0	0	1
1	0	0	0	0	1

Matrix Form of "A"

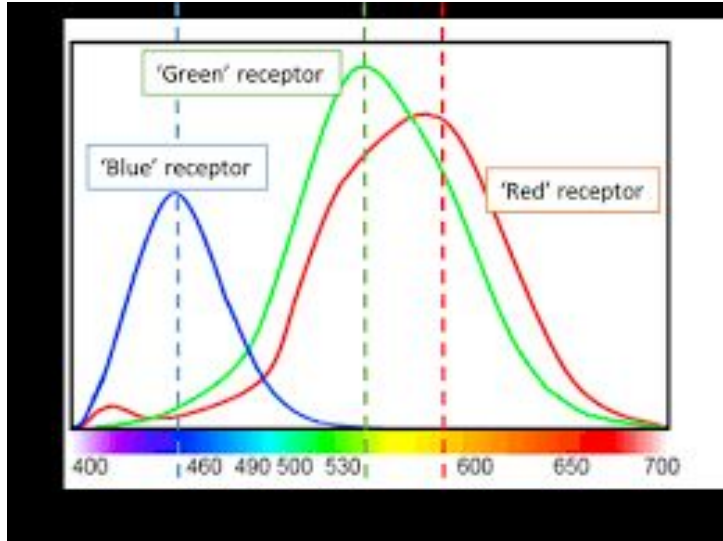
- Each pixel in the image corresponds to some value (i, j) in the matrix
- Values in the matrix represent intensity \Rightarrow ranges from 0 to 255 (normalized: 0 to 1)
 - 0 means black (no intensity)
 - 1 means white (maximum intensity)

Color vs. Grayscale Images




- Representing grayscale images is pretty easy!
 - Requires one **channel** to represent the image
 - Number of channels tells us how many numbers are required to specify each pixel's color/intensity
- Color images:
 - Requires 3 channels: Red, Green, and Blue (RGB)


Color Images



- Cones in the human eye (receptors) perceive light roughly in terms of Red, Green, Blue
- When representing color images in matrix form, we actually require three matrices for R, G, B intensities
 - Stacked together to create a single matrix representing the image where (i, j) in the matrix a tuple containing intensities for red, green, and blue




Function representation



- Images can be represented as continuous functions, where $f: (i, j) \rightarrow I$ (intensity)
- Function representation is **continuous**, as opposed to the matrix representation, which is **discrete**
 - This makes function representations useful for certain mathematical operations, such as calculating gradients
- For grayscale images, $f(x, y) = i$
- For color images, function f is **vector-valued** $f(x, y) = (f_r(x, y), f_g(x, y), f_b(x, y))$

Digital Images



- Use of sensor arrays
- Digital images are a discrete, 2D, digital representation of the function representation
 - Sampled and quantized

Image Transformations

What is a transformation? What can you with transformations?

Image Transformations

What is a transformation? What can you with transformations?

- Dimming/Brightening
- Rotating
- Mirroring
- much, much more!

⇒ Image Filtering


Image Transformation Operations

- Point
- Local
- Global

Point


Output value (m_i, n_j) is dependent upon the input value at the same coordinate.

- Pixel to pixel




Local

Output value (m_i, n_j) is dependent upon the input values of pixels in a $p \times p$ neighborhood of the corresponding pixel in the input image.



Global

Output value (m_i, n_j) is dependent upon all values in the input image.



Point Operations: Applications

- Dimming/brightening
- Mirroring
- Increasing/decreasing contrast
- etc.

Local Operations: Applications

- Reducing noise
- Filtering
- Convolutions
- etc.

Global Operations: Applications

- Fourier transforms
 - The Fourier transform breaks down an image into a sum of complex exponentials with different frequencies, phases, and magnitudes.
 - The output of the Fourier transform is a representation of the image in the frequency domain

Reducing Noise: What is noise?



image_A530



image_A733



image_C55



image_C70



image_CAS



image_KOD



image_NK710



image_OLM

Types of noise

- **Salt and pepper noise:** random occurrences of black and white pixels
- **Impulse noise:** random occurrences of white pixels
- **Gaussian noise:** variations in intensity drawn from a Gaussian normal distribution



Original



Salt and pepper noise





Impulse noise



Gaussian noise

Gaussian noise



$$f(x, y) = \overbrace{\hat{f}(x, y)}^{\text{Ideal Image}} + \overbrace{\eta(x, y)}^{\text{Noise process}}$$

Gaussian i.i.d. ("white") noise:
 $\eta(x, y) \sim \mathcal{N}(\mu, \sigma)$

```
>> noise = randn(size(im)).*sigma;  
>> output = im + noise;
```


Moving Average

- Take an $n \times n$ window and simply take the average of all the pixels in that window
- This becomes the value of the center pixel

Moving Average In 2D

$F[x, y]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$G[x, y]$

	0								

Moving Average In 2D

$F[x, y]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$G[x, y]$

	0	10							

Moving Average In 2D

$F[x, y]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$G[x, y]$

	0	10	20						

Moving Average In 2D

$$F[x, y]$$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$$G[x, y]$$

	0	10	20	30					

Moving Average In 2D

$$F[x, y]$$

[illegible]

$$G[x, y]$$

[illegible]

Moving Average In 2D

$$F[x, y]$$

[illegible]

$$G[x, y]$$

[illegible]

Moving Average In 2D

$$F[x, y]$$

[illegible]

$$G[x, y]$$

[illegible]

Image Filtering

- A linear filter is where an output pixel's value is a weighted sum of pixel values within a small neighborhood N
- The weights that you multiply each pixel by are called the kernel/filter/mask

Correlation filtering

$$G[i, j] = \underbrace{\frac{1}{(2k+1)^2}}_{\text{Attribute uniform weight to each pixel}} \underbrace{\sum_{u=-k}^k \sum_{v=-k}^k F[i+u, j+v]}_{\text{Loop over all pixels in neighborhood around image pixel } F[i,j]}$$

Cross-Correlation

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k \underbrace{H[u, v]}_{\text{Non-uniform weights}} F[i+u, j+v]$$

denoted $G = H \otimes F$

What can you do with correlation/cross-correlation filtering?

- Blur images (using a Gaussian filter)
- Shift images
- Dim/Brighten images
- Increase/Decrease contrast
- Edge detection
- etc.

Practice with linear filters



Original

0	0	0
0	1	0
0	0	0

?

Practice with linear filters



Original

0	0	0
0	1	0
0	0	0



Filtered
(no change)



Original

0	0	0
0	0	1
0	0	0

?



Original

0	0	0
0	0	1
0	0	0



Shifted left
by 1 pixel
with
correlation



Original

$\frac{1}{9}$

1	1	1
1	1	1
1	1	1


?



Original

$$\frac{1}{9}$$

1	1	1
1	1	1
1	1	1



Blur (with a
box filter)



Original

0	0	0
0	2	0
0	0	0

-

$\frac{1}{9}$

1	1	1
1	1	1
1	1	1

?




Original

0	0	0
0	2	0
0	0	0

- $\frac{1}{9}$

1	1	1
1	1	1
1	1	1



Sharpening filter:
accentuates differences
with local average

